

The  $5 \times 5$  matrix represents the distances between nodes. Each cell gives the length of a path of the node to the left of the cell to the nodes with the label that is shown on top of the table. The entry  $\infty$  indicates that the path length has not yet been computed.

The algorithm can now start!

If you hover the cursor above a cell, you can see the current path in the graph on the left.

The original matrix looks as follows:

	a	b	c	h	i
a	0	3	$\infty$	$\infty$	4
b	$\infty$	0	8	$\infty$	$\infty$
c	10	$\infty$	0	4	$\infty$
h	2	$\infty$	$\infty$	0	$\infty$
i	$\infty$	2	$\infty$	$\infty$	0

NB

$$\begin{cases} (i=0) \wedge (k=0 \vee j=0) \rightarrow \infty \\ (k=0) \wedge (i=0 \vee j=0) \rightarrow \infty \\ (j=0) \wedge (i=0 \vee k=0) \rightarrow \infty \end{cases}$$

K	0	1	2	3	4
i	0	1	2	3	4
j	0	1	2	3	4

$$\begin{aligned} m[i][k] + m[k][j] &< m[0][j] \\ m[2][0] + m[0][4] &< m[2][4] \\ 10 + 3 &< \infty \end{aligned}$$

	a	b	c	h	i
a	0	3	$\infty$	$\infty$	4
b	$\infty$	0	8	$\infty$	$\infty$
c	10	$\infty$	0	4	$\infty$
h	2	$\infty$	$\infty$	0	$\infty$
i	$\infty$	2	$\infty$	$\infty$	0

K	0	1	2	3	4
i	0	1	2	3	4
j	0	1	2	3	4

$$10 + 4 < \infty$$

	a	b	c	h	i
a	0	3	$\infty$	$\infty$	4
b	$\infty$	0	8	$\infty$	$\infty$
c	10	$\infty$	0	4	$\infty$
h	2	$\infty$	$\infty$	0	$\infty$
i	$\infty$	2	$\infty$	$\infty$	0

K	0	1	2	3	4
i	0	1	2	3	4
j	0	1	2	3	4

$$13 + 8 < \infty$$



gestire concorrenza questo valore non è garantito essere già presente in un ambiente parallelo!

Nota  $[i][j]$  non si riferisce mai se non quando  $K$  passa al valore successivo

"  $[k][j]$  " " " " "  $i$  " "

```

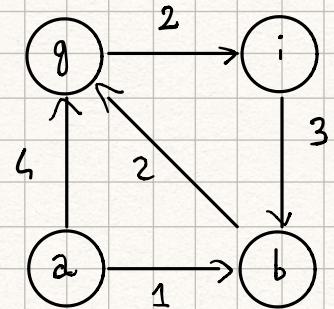
for (K=0; K<n; K++) {
    for (i=0; i<n; i++) {
        for (j=0; j<n; j++) {
            if (m[i][K] + m[K][j] < m[i][j]) {
                m[i][j] = m[i][K] + m[K][j];
            }
        }
    }
}

```

$K : J$

0 1 →  
0 →  
→  
→  
→  
→  
→  
→

	a	b	g	i
a	$\infty$	1	4	$\infty$
b	$\infty$	$\infty$	2	$\infty$
g	$\infty$	$\infty$	$\infty$	2
i	$\infty$	3	$\infty$	$\infty$



	a	b	g	i
a	$\infty$	1	3	5
b	$\infty$	$\infty$	2	4
g	$\infty$	5	$\infty$	2
i	$\infty$	3	5	$\infty$

$$\begin{aligned}
 [0][1] + [1][2] &< [0][2] \Rightarrow [0][2] = 3 \\
 1 + 2 &< 5 \\
 [0][2] + [2][3] &< \infty \\
 3 + 2 &< \infty
 \end{aligned}$$

↓

K : J	0 0	0 1	0 2	0 3	0 4	0 5	0 6	0 7	0 8	0 9
0 0	0	1	X							
0 1		0	2	X						
0 2			0	3	X					
0 3				0	4	X				
0 4					0	5	X			
0 5						0	6	X		
0 6							0	7	X	
0 7								0	8	X
0 8									0	9
0 9										0

	a	b	g	i
a	$\infty$	1	4	$\infty$
b	$\infty$	$\infty$	2	$\infty$
g	$\infty$	$\infty$	$\infty$	2
i	$\infty$	3	$\infty$	$\infty$

$k$  row col  
1  $t_h[0][0] b[0][1] \rightarrow [0][2] \checkmark$   
 $k$   $< 4$   
2  $t_h[0][1] b[0][1] \rightarrow [0][3] \checkmark$   
 $< \infty$

Se  $k$  è un indice che permette di tenere traccia del modo intermedio.

1	0	1
1	0	2
1	0	3
1	1	0

$$\begin{array}{c} 1 + \infty \\ 1 + 2 \\ 1 + \infty \\ \infty + \end{array} < 4 \quad \checkmark \quad \times$$

$$[0][2] = 3$$

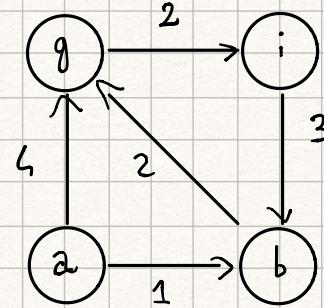
1	1	2
1	1	3
1	2	0
1	2	1

1	2	3
1	3	0
1	3	1
1	3	2

2	0	0
2	0	1
2	0	2
2	0	3
2	1	0

2	1	2
2	1	3
2	2	0
2	2	1
2	2	3
2	3	0
2	3	1
2	3	2

	a	b	g	i
a	$\infty$	1	4	$\infty$
b	$\infty$	$\infty$	2	$\infty$
g	$\infty$	$\infty$	$\infty$	2
i	$\infty$	3	$\infty$	$\infty$



$$b \rightarrow g \Rightarrow 2$$

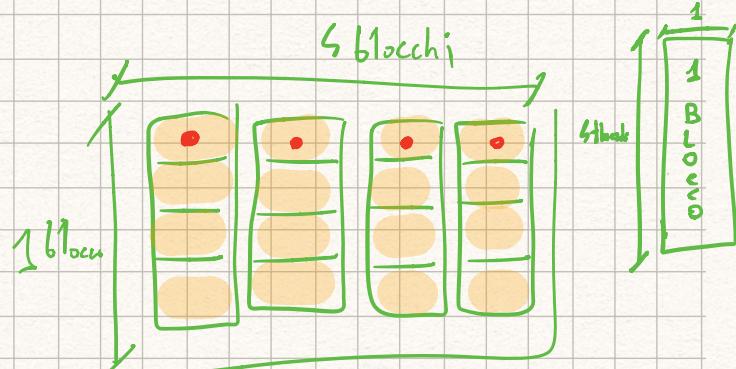
$$g \rightarrow b \Rightarrow \infty$$

- Ad ogni iterazione il  $K$  è lo stesso per tutte le thread e per tutti i blocchi
- Ogni volta che cambia la  $i$  si ripetono gli stessi  $[K][j]$

possibile usare memoria condivisa

3	0	1
3	0	2
3	0	3
3	1	0
3	1	2
3	1	3
3	2	0
3	2	1
3	2	3
3	3	0
3	3	1
3	3	2

	K	i	j
0	0	0	0
0	0	1	X
0	0	2	X
0	0	3	X
0	1	0	X
0	1	1	X
0	1	2	X
0	1	3	X
0	2	0	X
0	2	1	X
0	2	2	X
0	2	3	X
0	3	0	X
0	3	1	X
0	3	2	X
0	3	3	X
0	3	4	X



Ogni blocco ha le seguenti thread:  
 $th[0][0] \rightarrow K-j\_value = N[2^*n+2] =$   
 $th[0][1]$   
 $th[0][2]$   
 $th[0][3]$

## BLOCCI      THREADS

x	y
[0][0]	[0][0]
[1][0]	[0][1]
[2][0]	[0][2]
[3][0]	[0][3]

x	y
[0][0]	[0][0]
[1][0]	[0][1]
[2][0]	[0][2]
[3][0]	[0][3]

i	j
0	0
1	0
2	0
3	0
0	1
1	1
2	1
3	1
0	2
1	2
2	2
3	2
0	3
1	3
2	3
3	3

Valori calcolati dal blocco 0

Valori calcolati dal blocco 1

Valori calcolati dal blocco 2

Valori calcolati dal blocco 3

$$tx + ty * \text{BlockDim} / \text{warpSize} = \text{warpID}$$

$$0 + 0 / 32 = 0$$

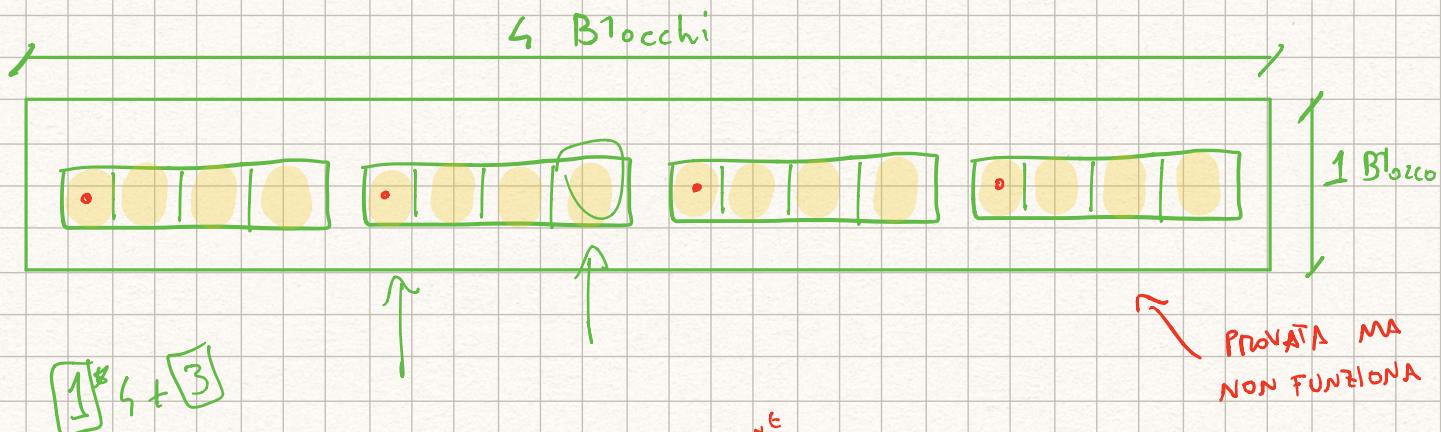
$$0 + 4 / 32$$

$$0 + 8 / 32$$

$$0 + 12 / 32$$



$\Rightarrow$  WARP DIVERGENCE + ACCESSI A MEMORIA NON COALIZZATI



## BLOCCI      THREADS

x	y
[0][0]	[0][0]
[0][1]	[1][0]
[0][2]	[2][0]
[0][3]	[3][0]

x	y
[0][0]	[0][0]
[1][0]	[0][1]
[2][0]	[0][2]
[3][0]	[0][3]

i	j
0	0
0	1
0	2
0	3
1	0
1	1
1	2
1	3
2	0

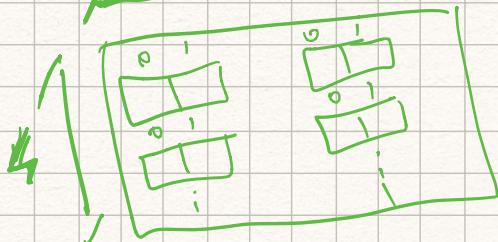
Valori calcolati dal blocco 0

Valori calcolati dal blocco 1

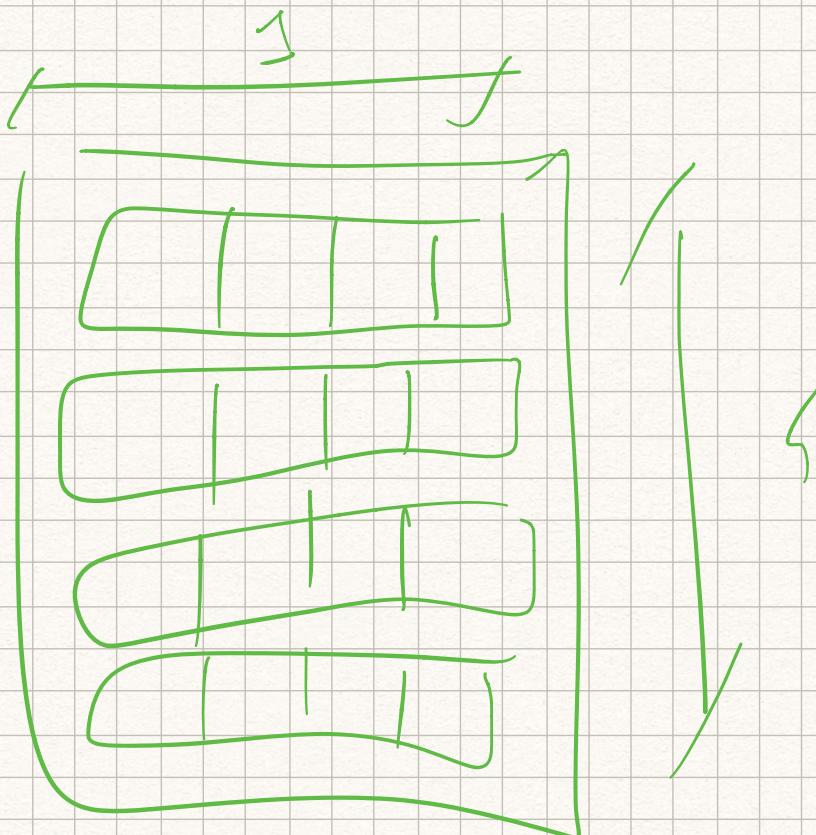
NON SI PUÒ FAR E PERCHÉ  
LA "i" NON VARIA NELLO STESSO  
Blocco, QUINDI GLI STESSI [i][j]  
SI RIPETONO SOLO IN BLOCCI DIVERSI

MA AVRÀ UNI  
STESSI [i][j]  
AVENDO COME  
VARIANTE NELL  
STESO BLOCCO  
LA "j"





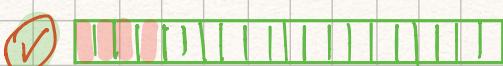
2	1		Valori cattolati dal blocco 2
2	2		
2	3		
3	0		
3	1		Valori cattolati dal blocco 3
3	2		
3	3		



i	j	
0	0	blocco 0
1	0	
2	0	
3	0	
0	1	
1	1	blocco 1
2	1	
3	1	

$$tx + ty * \text{BlockDim} / \text{WarpSize} = \text{WarpID}$$

$$\begin{array}{ll} 0 + 0 & / 32 \\ 1 + 0 & / 32 \\ 2 + 0 & / 32 \\ 3 + 0 & / 32 \end{array}$$



PROBLEMA

RIMANGONO GLI ACCESSI A MEMORIA NON CODIZZATI.

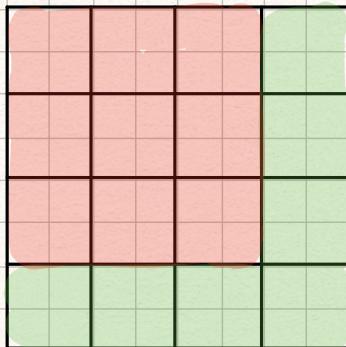
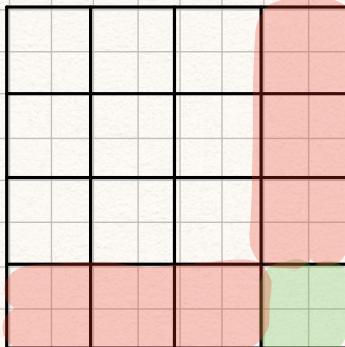
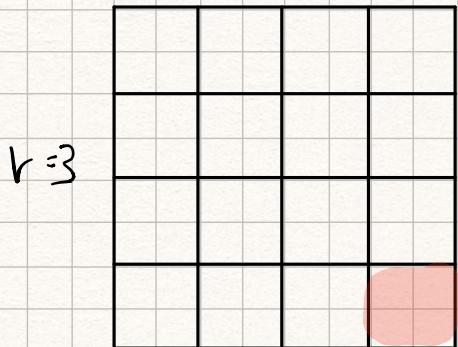
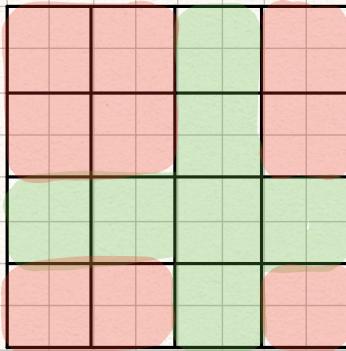
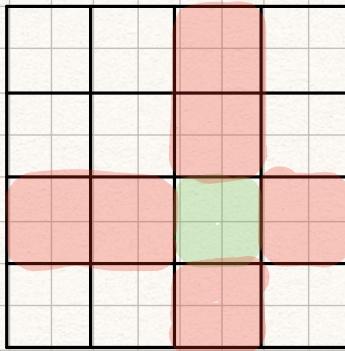
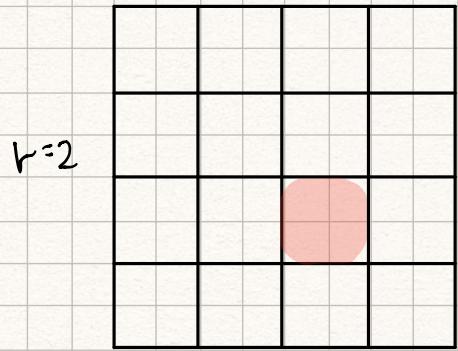
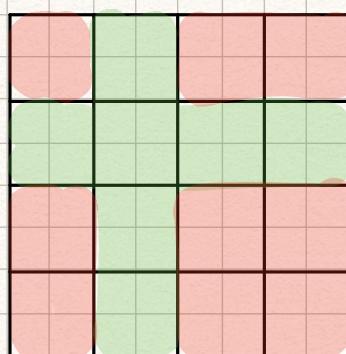
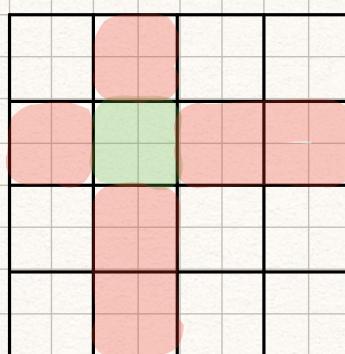
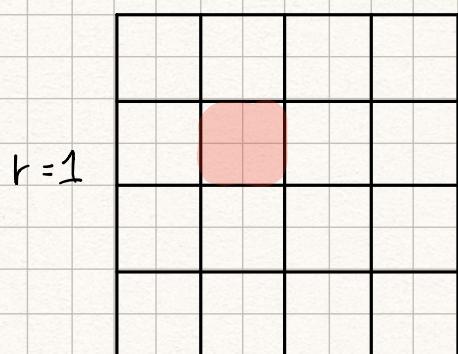
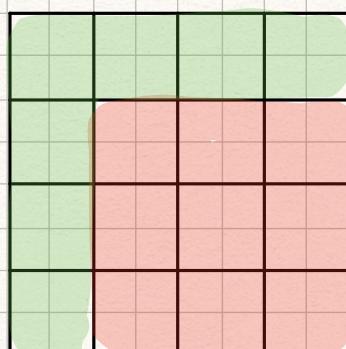
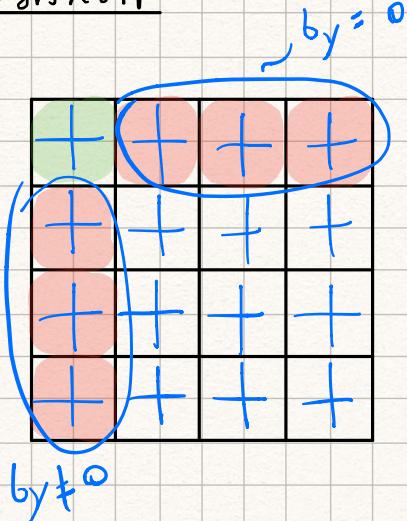
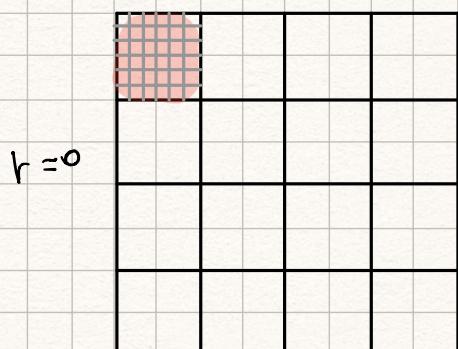
Notiamo quindi che con questo approccio, nonostante l'uso di memoria condizionata non si hanno grandi risultati.

ph2,0      ph 2,1

$th^{(0,0)} 0,2 \Rightarrow 1,0$   
 $th^{(0,1)} 0,3 \Rightarrow 1,1$

# Blocked Floyd Warshall

$th(i,j)$        $0, i$   
 $i, n-1$   
 $i, j$   
 $i, n-1$



Esempio

K = 1

← stage

Phase 1:  $D^o(i,j) = A^o(i,j) = cost(i,j)$

$$D^k(i,j) = A^k(i,j) = \min \{ A^{k-1}(i,j), A^{k-1}(i,k) + A^{k-1}(k,j) \}$$

$$\begin{aligned} d^{(1)}(0,0) &= \min(\infty, 1 + \infty) = \infty \\ d^{(1)}(0,1) &= \min(1, 1 + \infty) = 1 \\ d^{(1)}(1,0) &= \min(\infty, \infty + \infty) = \infty \\ d^{(1)}(1,1) &= \min(\infty, \infty + \infty) = \infty \end{aligned}$$

	a	b	g	i
a	$\infty$	1	4	$\infty$
b	$\infty$	$\infty$	2	$\infty$
g	$\infty$	$\infty$	$\infty$	2
i	$\infty$	3	$\infty$	$\infty$

Phase 2:  $D^o(i,j) = A^o(i,j)$

$$D^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \} \text{ rows}$$

$$D^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \} \text{ cols}$$

$$\begin{aligned} d^{(1)}(0,2) &= \min(4, 1 + 2) = 3 \\ d^{(1)}(0,3) &= \min(\infty, 1 + \infty) = \infty \\ d^{(1)}(1,2) &= \min(2, \infty + 2) = 2 \\ d^{(1)}(1,3) &= \min(\infty, \infty + \infty) = \infty \end{aligned}$$

rows

	a	b	g	i
a	$\infty$	1	3	$\infty$
b	$\infty$	$\infty$	2	$\infty$
g	$\infty$	$\infty$	$\infty$	2
i	$\infty$	3	$\infty$	$\infty$

$$\begin{aligned} d^{(1)}(2,0) &= \min(\infty, \infty + \infty) = \infty \\ d^{(1)}(2,1) &= \min(\infty, \infty + \infty) = \infty \\ d^{(1)}(3,0) &= \min(\infty, 3 + \infty) = \infty \\ d^{(1)}(3,1) &= \min(3, 3 + \infty) = 3 \end{aligned}$$

cols

Phase 3:  $D^o(i,j) = A^o(i,j)$

$$D^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \}$$

$$\begin{aligned} d^{(1)}(2,2) &= \min(\infty, \infty + 2) = \infty \\ d^{(1)}(3,3) &= \min(2, \infty + \infty) = 2 \\ d^{(1)}(3,2) &= \min(\infty, 3 + 2) = 5 \\ d^{(1)}(3,3) &= \min(\infty, 3 + \infty) = \infty \end{aligned}$$

$\infty$	1	3	$\infty$
$\infty$	$\infty$	2	$\infty$
$\infty$	$\infty$	$\infty$	2
$\infty$	3	$\infty$	5

$$\begin{array}{ll} PRIM\_R & PRIM\_C \\ (2,2) & (0,0) + (0,0) \\ (2,2) & (0,1) + (1,0) \end{array}$$

$$\begin{array}{l} \infty + 3 \\ \infty + 2 \end{array}$$

K = 2

Phase 1

$$D^o(i,j) = A^o(i,j) = \text{cost}(i,j)$$

$$D^k(i,j) = A^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \}$$

duplica su questi indici!!!

$$\begin{aligned} d^{(2)}(2,2) &= \min \{ \infty, 3 + \infty \} = \infty \\ d^{(2)}(2,3) &= \min \{ 2, 3 + \infty \} = 2 \\ d^{(2)}(3,2) &= \min \{ 5, 2 + \infty \} = 5 \\ d^{(2)}(3,3) &= \min \{ \infty, 2 + \infty \} = \infty \end{aligned}$$

	a	b	g	i
a	$\infty$	1	3	$\infty$
b	$\infty$	$\infty$	2	$\infty$
g	$\infty$	$\infty$	$\infty$	2
i	$\infty$	3	5	$\infty$

Phase 2

$$D^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \} \text{ rows}$$

$$D^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \} \text{ cols}$$

$$\begin{aligned} d^{(2)}(2,0) &= \min \{ \infty, \infty + \infty \} = \infty \\ d^{(2)}(2,1) &= \min \{ \infty, 2 + 3 \} = 5 \quad \text{rows} \\ d^{(2)}(3,0) &= \min \{ \infty, 5 + \infty \} = \infty \\ d^{(2)}(3,1) &= \min \{ 3, 5 + \infty \} = 3 \end{aligned}$$

	a	b	g	i
a	$\infty$	1	3	$\frac{5}{\cancel{\infty}}$
b	$\infty$	$\infty$	2	$\frac{5}{\cancel{\infty}}$
g	$\infty$	$\cancel{\infty}$	$\infty$	2
i	$\infty$	3	5	green

$$\begin{aligned} d^{(2)}(0,2) &= \min \{ 3, 3 + \infty \} = 3 \\ d^{(2)}(0,3) &= \min \{ \infty, 3 + 2 \} = 5 \\ d^{(2)}(1,2) &= \min \{ 2, 2 + \infty \} = 2 \\ d^{(2)}(1,3) &= \min \{ \infty, 2 + 2 \} = 4 \end{aligned}$$

cols

$(2,1), (0,0) + (0,1) \rightarrow + \infty$   
 $(2,1), (0,1) + (1,1) \rightarrow 2 + 3$

Phase 3

$$D^k(i,j) = \min \{ D^{k-1}(i,j), D^{k-1}(i,k) + D^{k-1}(k,j) \}$$

$$\begin{aligned} d^{(2)}(0,0) &= \min \{ \infty, 3 + \infty \} = \infty \\ d^{(2)}(0,1) &= \min \{ 1, 3 + \infty \} = 1 \\ d^{(2)}(1,0) &= \min \{ \infty, 2 + \infty \} = \infty \\ d^{(2)}(1,1) &= \min \{ \infty, 2 + \infty \} = \infty \end{aligned}$$

$$(1,1) \quad (1,0) + (0,1) \quad \infty + 1$$

$$(1,1) \quad (1,2) + (2,1) \quad 2 + 1$$

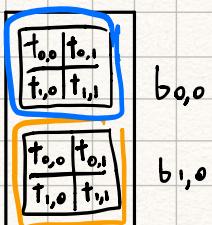
	a	b	g	i
a	$\infty$	1	3	5
b	$\infty$	$\infty$	2	4
g	$\infty$	5	$\infty$	2
i	$\infty$	3	5	$\infty$



STAGES = 2

$$0 \leq b_x < \text{stages} - 1$$

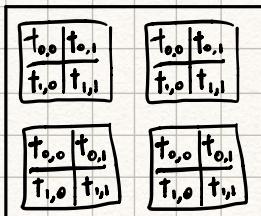
PHASE 2  $\rightarrow$  Kernel config.



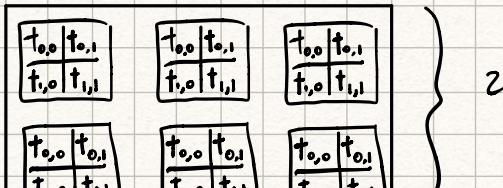
<<< stages - 1, 2 >>>

stage = K

$$0 \leq K < \text{stages}$$



i blocchi di questo rego  
conterranno le righe di N  $\rightarrow$



i blocchi di questo rego  $\rightarrow$   
conterranno le colonne di N

stages - 1

○	○	○	○	○	○	○
○	○	○	○	○	○	○
○	○	○	○	○	○	○
○	○	○	○	○	○	○
○	○	○	○	○	○	○

8 8 8 8 8 8 8 8 8 8 |

808 9900  
9907

0 1 2 3 4 5 6 7 8 - - - -

100  
200  
300  
400  
500  
600  
700

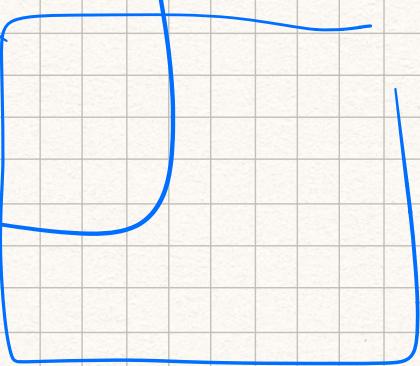
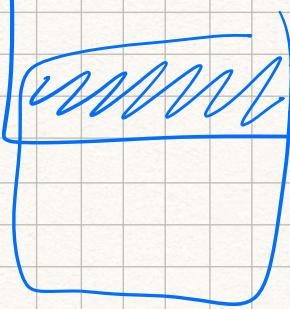
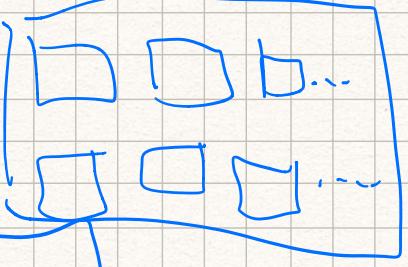
8 x 8

808

8 x 8

99

1



# Progetto Architetture Avanzate A.A. 2017-2018

Parallelizzazione algoritmo All Pairs Shortest Path (Floyd-Warshall) in CUDA.

- ) Codice sorgente sequenziale C/C++ consegnato dal docente (e-learning).
- ) Benchmark di riferimento per i test consegnati da docente (e-learning).

## Obiettivo 1:

- ✓ 1) Analisi codice sequenziale con profiler
- ✓ 2) Sviluppo kernel CUDA per esecuzione su GPU con ottimizzazione su
  - a. Kernel configuration // provare varie e tenere la migliore
  - b. Branch divergence
  - c. Memory coalescence
  - d. Shared memory usage
  - e. Load balancing
  - f. Tecniche avanzate (ILP, etc.)

## Risultati sperimentali:

Speedup rispetto sequenziale ottenuto con l'applicazione incrementale delle ottimizzazioni, con presentazione risultati profiler CUDA ad ogni step. Risultati (tempi) ottenuti su tutti i benchmark.

1) Dopo una prima analisi con Valgrind è possibile notare che il codice di lavoro maggiore si ha all'interno della funzione che implementa l'algoritmo Floyd-Warshall.

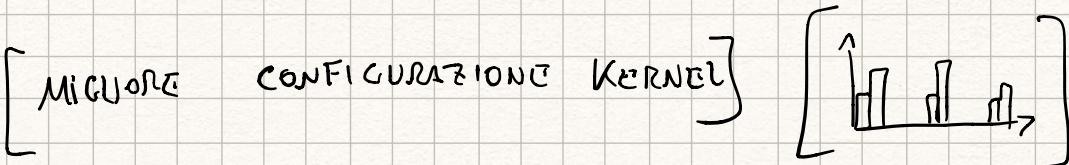
In particolare su tutti i benchmark il codice si concentra sulle istruzioni condizionali all'interno del ciclo più interno!

Dopo una fine accortezza il codice è stato leggermente ottimizzato risformando  $n$  iterazioni per tutte le volte che la matrice[i][k]

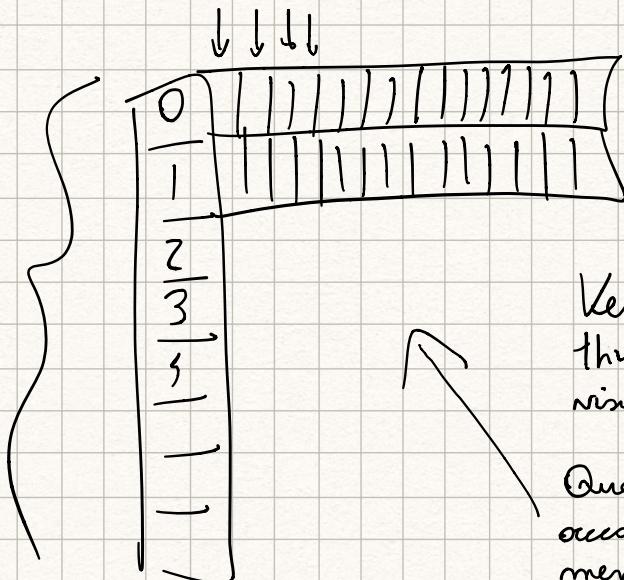
ha un valore pari ad INF, e lasciando inviato il gergo tutte le volte che  $i = j$ , ossia tutte le volte che si sta per computare il commento di costo minimo da un modo a se stesso.

2) La prima versione parallela del Floyd-Warshall vuole focalizzarsi sulla funzionalità dell'algoritmo, senza troppo importanza alle performance.

Provando questa versione Naive sui nostri benchmark, si può già notare un buon miglioramento in termini di performance, in particolare mi ha un migliore rendimento quando il Kernel presenta la seguente configurazione.

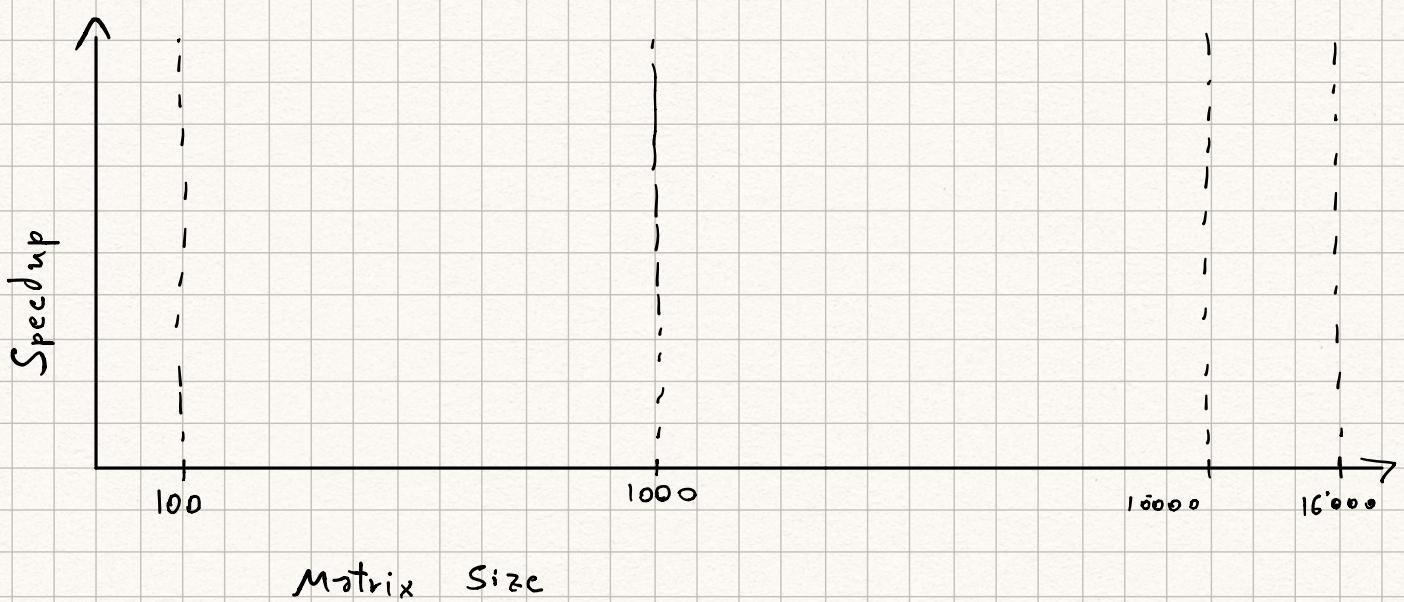
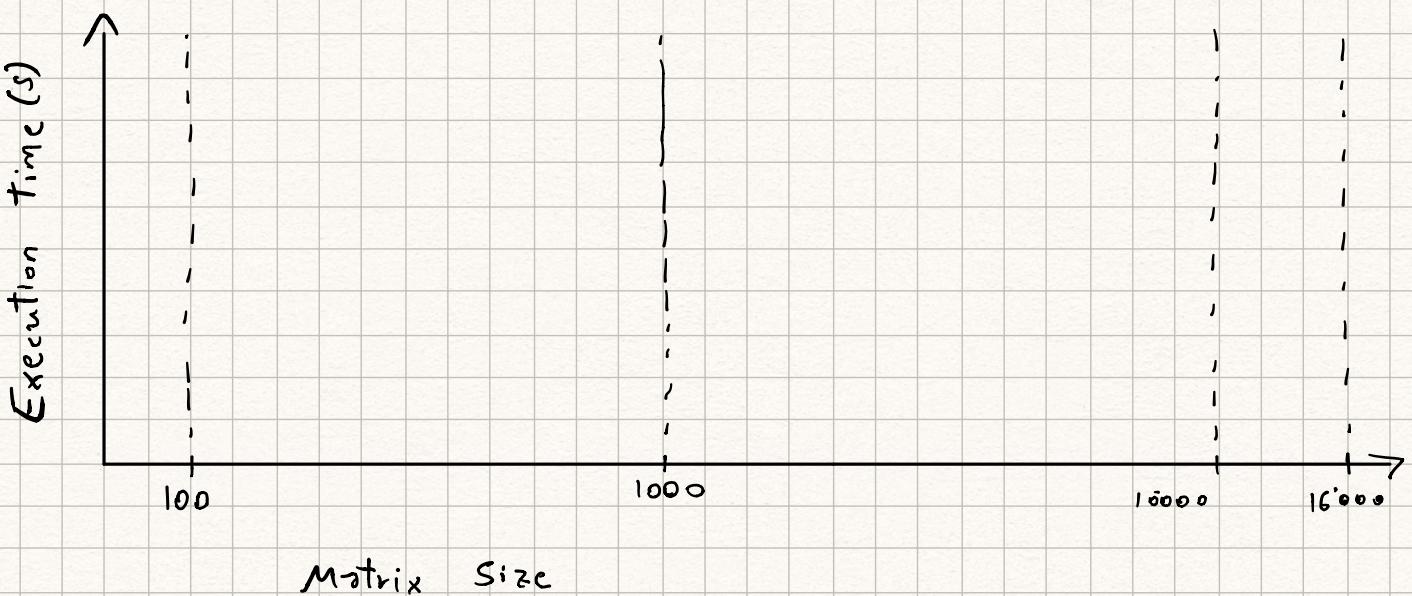


I PRIMI 32 WARPS  
FORMANO IL  
PRIMO BLOCCO



Verificare numeri e quali thread sono attive con visual studio!

Questi perché le thread occorrono sequenzialmente in memoria, ma non è detto che vengono eseguite in maniera sequenziale (thread dello stesso warp)

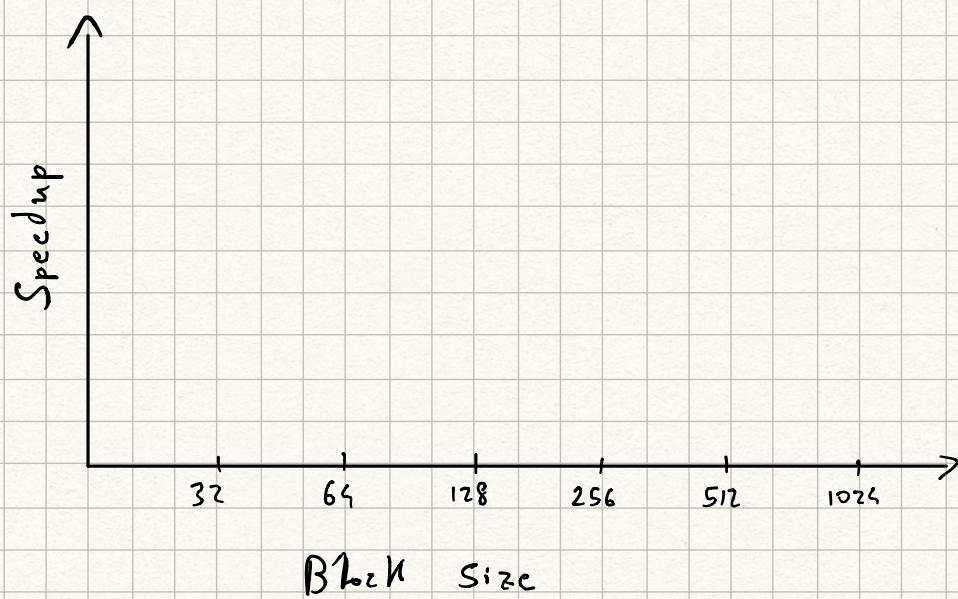


Questi 2 grafici vorrei belli da mettere nei risultati finali.



// Questi grafici  
vanno bene da  
mettere nella sezione  
CUDA.

Prendiamo come  
riferimento il  
 $C-64 = 10.000$  modi



th	time	speedup
32	74 ms	17x
64	52 ms	25x
128	52 ms	25x
256	53 ms	26x
512	53 ms	25x
1024	57 ms	23x

1000 nodi

$$S = 10$$

$$10 + 1$$

$$11$$

$$10 + 2$$

$$12$$

$$10 + 3$$

$$13$$

$$10 + 4$$

$$14$$

$$S = 10$$

$$S = 10 + 1$$

$$11$$

$$S = 11 + 2$$

$$13$$